

Hardware Determinism over Statistical Gradients:

A GPU-Accelerated Boolean Attack on the Final U-534
Enigma M4 Ciphertext (P1030680)

Digital Defiance
HELUT Campaign Report

August 2026

Abstract

For decades, message P1030680 has remained the sole unbroken Enigma M4 transmission from the May 1, 1945 intercept of German submarine U-534. Modern attempts to decrypt this 72-letter ciphertext have historically relied on statistical hill-climbing and distributed computing. This report shows that purely statistical attacks fail against short M4 ciphertexts due to plugboard overfitting, documents a mid-message middle-rotor turnover flaw in whole-message trigram scoring, and describes a software Welchman diagonal board on Apple Silicon Metal (40 million machine settings per second) with a mandatory ≤ 10 -plug completion kill chain. We record graded controls (known-message break), archival eliminations, clean negatives under catalog and operational cribs, a quarantine path for soft near-misses, and a parallel TensorLUT continuous-discrete compiler track. A runtime benchmark selected a four-lane unified Verilog Future Bank. The unified P1030680 Mulein search machine and overall campaign carry the formal name *Nazi Blaster 9000*; its historical 261-entry canonical-menu/Future campaign is *Fahrenheit 261*. That historical name denotes the inventory and campaign, not executed breadth: its bounded launch ran only identity Future 0 at settings $0 \leq s < 256$, leaving 260 Futures unexecuted. A separate 628-Future operational-prior preflight then covered setting 0 across all identity and post-gap entries: 16,328 checked receipts produced 15 host-replayed one-edge repaired physical candidates, all non-exact, with zero exact hits and zero BREAK gates. The same shell and all 628 Futures then completed settings $1 \leq s < 256$: 10,048 synchronized chunks and 4,163,640 checked receipt projections retained 4,184 host-verified physical candidates but fired zero BREAK gates. Forty exact identity hits died at ≤ 10 -plug completion; repaired identity and all post-gap candidates still require correction or geometry-aware full-message replay, and none supplied a scored plaintext, key, or decrypt. This is bounded evidence, not global coverage or evidence that the ciphertext is garbled. **We do not claim a decrypt of P1030680.** Campaign fitness is cleartext Metal batch, not HELUT's encrypted torus path (documented separately in `paper/helut.tex`).

1 The Failure of Statistical Gradients on Short Ciphertexts

When attempting to break an Enigma ciphertext, modern cryptographic approaches default to statistical optimization, such as Index of Coincidence (IC), n -gram fitness, and genetic algorithms. However, empirical testing against P1030680 proves that statistics cannot *search* a 72-letter Enigma space; they can only *rank* a small list of survivors.

The Enigma plugboard (the *Steckerbrett*) offers roughly 47 bits of freedom. On a short message, this immense freedom causes statistical models to overfit. Algorithms will output high-scoring strings of letters that closely resemble German syntax, but these strings are mathematical hallucinations. Control tests on a known, solved message (P1030684) confirmed that purely statistical engines favor these false positives over the true key. Consequently, any successful attack on short M4 intercepts must abandon statistical guessing in favor of absolute Boolean logic.

2 Emulating the Welchman Diagonal Board in GPU Shaders

To execute a deterministic attack, the hardware architecture of the 1945 Enigma must be perfectly mapped to modern hardware. Using Homomorphic Edge Look-Up Tensors (HELUT), we engineered a digital Bombe inside Apple Silicon Metal shaders.

This software-defined FPGA evaluates Gordon Welchman’s diagonal board graph reductions at a rate of 40 million machine settings per second. The engine utilizes Turing-shaped reductions based on the physical reality of the short-message M4:

- Because the far-left Enigma rotor (the Greek wheel) never steps during a short message, its ring can be digitally pinned.
- Because the left rotor’s notch drives nothing, its ring can also be pinned.
- Together, this physical constraint collapses the total search space by a massive factor of 676.

The mathematical soundness of this reduction was verified in a blind control run. Given a 27-letter crib from a known message (P1030684) and no other parameters, the GPU pipeline evaluated 16 billion settings and extracted the true key, exact shell, and all ten historical plugboard cables in just 361 seconds (claim C2). The archived log (`logs/control-p1030684-rings.log`) reads 427 s: 361 s was the first run, and the log was overwritten by a re-run after head scoring was added. Both figures are recorded in `BREAK.P1030680.md`; the re-run is the reproducible one.

3 Physical Constraints vs. Mathematical Ghosts

Translating mathematical graph theory into physical 1945 hardware exposed a significant vulnerability in pure Boolean reductions. Certain crib wedges produce valid Welchman states that are mathematically flawless but physically impossible.

We classify these false positives as “ghosts”. In a trial run, a crib of `UUUVIRSIBENNULEINS` at offset 0 produced 12 mathematical survivors out of the Welchman board. However, when mapping the remaining letters, none of these 12 settings could be completed within the physical limit of exactly 10 plugboard cables. Forensics revealed these were “split menus,” where connected components within the graph admitted zero valid starting seeds.

To prevent these ghosts from overwhelming the CPU, we built a SAT-completion kill chain inline as the GPU drains. This ensures that only mathematically valid states that *also* fit physical hardware constraints survive the hardware pipeline (claim C3).

4 The Mid-Message Turnover Scoring Flaw

Our research uncovered a systemic flaw in how distributed computing projects historically score Enigma decrypts, likely explaining the decades-long failure to crack P1030680.

Standard algorithmic sweeps utilize a linguistic scanner to score the full 72-letter plaintext. However, if the true key clicks the middle rotor over *after* the crib phrase but *before* the end of the message, the plaintext will be perfectly decrypted across the crib span, but turn to garbage for the remainder of the message.

Because traditional trigram discriminators score the *entire* message, the garbled tail drags the overall score down to the midpoint between German and random noise. This causes the engine to discard the true key, actively rejecting the correct answer.

To rectify this, we designed a **windowed discriminator score**. This metric evaluates the readable head of a decrypt (requiring at least 16 non-crib letters) independently from the whole message, passing the key if either the full message or the head meets the -3.600 threshold. Applying this fix revealed 2,219 physically valid ten-plug boards that had been previously rejected by the flawed whole-message scoring (claim C12).

5 The M-Thetis / Potsdam Register Mismatch

Cryptanalysis is limited by the quality of the plaintext hypotheses (cribs) used as wedges. Archival auditing of the May 1, 1945 U-534 corpus highlighted a severe historical blindspot regarding the target message.

The scraped corpus contains 48 broken M4 messages. Crucially, 31 originate from the Potsdam net, 16 from a second net, and 1 from a third. P1030680 is the only message in the scrape with the working assignment M-Thetis. That assignment remains supported by ACH under the later allocation list used by Hörenberg, but the catalogue chronology is not fully settled: Selm Merel Wenselaers reports an earlier list, *Forelle*, which maps column 645 to Aegir while the later list maps it to Thetis. Changing Zuteilungslisten can reconcile the apparent contradiction without changing the K-Buch column; the exact date, codename, and provenance of the later list remain pending. Likewise, preservation among U-534’s radio papers establishes recovery provenance, not intended recipient. An intercepted training-network signal is an open hypothesis, not a finding.

Because Bletchley Park never bothered to work the M-Thetis training net, and no other Thetis messages survive in the corpus, all 100 historical cribs used to attack P1030680 were register language borrowed from other operational networks. If Thetis operators opened their 72-letter messages with callsigns or key groups before the body—unlike Potsdam operators—the exact crib placements used in modern attacks are fundamentally the wrong shape.

That reading is correct about *provenance* and has been applied too broadly, which is worth stating because it changes where compute should go. Classifying the 868 tested placements of length ≥ 16 by content rather than by source, **736** contain only universal Kriegsmarine procedure—spelled numerals, VVV, UUU, FLOTTX, VONVON—and just **132** carry net-specific addressing such as KOMXADM, CHEF, or place names. Spelled numerals and procedure markers are service-wide conventions, not Potsdam idiolect, and the most connected of them concern *U-boat flotilla movements in the Baltic*, which is precisely what a Baltic training net existed to coordinate. The register objection therefore bites hard on openings and addressing and much more weakly on numerals and procedure. **288** of those universal-only placements are single-component with ≥ 8 loops and have never received middle-ring coverage, so a plausible-crib hypothesis remains live in *untested key space* rather than eliminated. The resilience of P1030680 is plausibly a mixture of mismatched operational doctrine and incomplete ring coverage, rather than the cryptographic strength of the cipher itself.

The archival side of that question is now active team work. Selm Merel Wenselaers, a historian and curator in Amsterdam and Antwerp, leads the campaign’s historical/archival arm alongside HELUT’s mechanical cryptanalysis. Her named DEFE and U-3521 KTB findings report a concrete process in which scuttling parties remained aboard while crews transferred to Neustadt; NEUSTADT and VERSENKUNG are therefore linked stages rather than competing scenarios. She also reports scheduled FT-Programmzeit reception during AGRU-Front training and has enquiries open with archives and collections for radio papers, key material, and a second Thetis message. These reports move priors but remain source-reported until scans or exact transcriptions enter the campaign receipts.

One associated radio finding is independently reproducible in the repository: P1030668 contains FFFTTTBLEIBTBESETZT exactly, while P1030707 carries the corrupted FFFTTTBLEIBTBESE0ZTX. The clean phrase therefore does not depend on reconstructing the damaged sister copy. Its 19-letter core is a weak menu, but the exact 40-letter P1030668 source context yields four single-component, 18-loop alignment hypotheses. They are staged, not run, and remain behind the full middle-ring arm. The clean U-3024/U-2538 number strings are withdrawn after raw-copy rechecking, and short historical terms are kept independent rather than concatenated into invented sentences.

Before the next GPU tranche begins, the campaign reports the completed findings and

exact proposed scope back to Selm and incorporates any correction or archive reply first. The findings are joint work: historical research defines the source-grounded questions, and the graded machinery turns them into bounded measurements.

Potsdam and Plaipe daily keys for 1 May were eliminated by exhaustion against P1030680 (claim C11): all 26^4 message keys land at the noise floor.

6 Graded Campaign Status (August 2026)

The operational ledger is `BREAK_P1030680.md`; the public claim freeze is `directives/claim-sheet.md`. Table 1 summarizes what is *proven*, what is *eliminated*, and what remains open. **None of these rows is a decrypt of P1030680** (non-claim N5).

Table 1: Campaign grades (cleartext Metal). Fitness is not HELUT encrypted tick rate (N6).

Result	Status
Welchman blind control (P1030684)	Break in 361 s first run / 427 s archived re-run; all 10 plugs
Stochastic KPA datapath	Control oracle + near-miss PASS; blind stecker $\sim 22/72$
Training / collapse / weather / RIGA priors	0 survivors under 80% rigor (ceiling $\sim 60\text{--}69\%$ = coincidence)
Exact ≥ 16 catalog $\times$ rings AAAA	Clean negative
Curated exact / fuzzed top-40 $\times$ right rings	Clean negative
UEBUNG / Thetis-register arms	Clean negative (register rings parked at 13/73)
Soft-tail UEBUNG quarantine escalate	55 seeds; Hybrid 0 survivors
Regenbogen / Hannibal legacy Welchman slices	Not BREAK where tested under AAAA/right-ring scope; soft escalate 0 survivors. This does not close correction-aware operational Futures
Own-orders / filler $\times$ AAAA	Own: 0 physical; filler below soft — no rings
Catalog exact $\times$ right rings	Parked originalIndex 417/2513; <code>--bombe-from 418</code>
Crib catalog is recurrence-filtered	Known but previously unrun surface. An internal 2026-08-16 relay audit measured 152,938 of 154,469 distinct 16–40 letter corpus windows (99.0%) as one-carrier phrases and expanded them to 2,691,455 legal placements; its 24-placement fixture was never run or printed publicly. Selm Merel Wenselaers (Historian / Curator, Amsterdam / Antwerp) independently flagged the filter and supplied three exact historical leads on 2026-09-22
Selm-led historical/cryptanalytic arm	Joint campaign work. Selm Merel Wenselaers leads the active archival arm; her source research and menu judgement define evidence-backed questions, and HELUT’s graded path turns them into controlled measurements. Exact and post-gap tests of her first three leads are both clean 28/28 negatives. A new exact P1030668 BLEIBTBESETZT context fixture is staged but not run; archive replies and the full-middle arm come first

Selm-team hapax leads $\times$ right rings	Clean negative. 28/28 placements dead at the board, 0 raw stops over 4.471×10^{11} settings at 38.0M/s, 0 physical, 0 quarantine. A verdict on 28 placements of the <i>recorded</i> ciphertext with the middle ring pinned A — not on the phrases generally. All three are second-net/Potsdam, so the imported-register objection stands
Same leads, post-gap $\delta=4$ geometry	Clean negative. 28/28 dead, 0 raw stops over 4.471×10^{11} settings at 37.8M/s, 0 physical, 0 quarantine — an identical settings count to the exact arm, since splicing preserves menu shape and moves only rotor coordinates. Both readings of the transcript are closed for these placements
Tolerance 1 on those menus	Measured affordable, not yet run. All 24 qualifying menus project 0 stops at a measured $7.7 \times$ (~ 19.7 h). That corrects the expectation that low-loop menus would inflate: a menu contradicting everywhere at tolerance 0 leaves nothing for tolerance to widen. With exact and post-gap both empty, a mis-read <i>letter</i> is the only edit hypothesis left for these phrases
Middle ring $\neq$ A	Partial / suspended at 1/24 against P1030680. The completed placement is dead over 4.152×10^{11} settings; resume <code>--bombe-menus 0 --bombe-middle-ring --bombe-ring-sweep --bombe-from 2</code> . The control breaks at full middle $\times$ right coverage ($26 \times$ space for $12.2 \times$ time)
Long menus (len 28–40)	Two maximum-contamination re-runs complete: Phase 16 body arm 40/40 dead with 0 raw stops; Phase 8 top-30 reproduces 73 raw stops with zero additional survivors and no break. The 300-place catalog remainder is retained but deprioritised; ≤ 25 letters unaffected
Ciphertext garble risk	Sister P1030681 needed degarbling; exact bars may drop true shells. 17 of 48 published U-534 keys fail to reproduce their own recorded plaintext, but from at least three causes — multi-part per-part keys, indels, and only 4 controls / 8 events of isolated substitution. An earlier “16 controls prove archive garble” reading is withdrawn
Crib-free (Ostwald/Weierud)	Still negative at 72 letters, but current-model exhaustion is now measured. Unseeded generic: 5/30, median -0.2306 . Single-plug exhaust-6 reaches -0.0672 (about 71 % of the deficit removed) and saturates by exhaust-10 at -0.0567 . New two-plug exhaust-6 gives 30 % wins / -0.0361 ; widening to exhaust-10 gives 40 % but worsens the median to -0.0486 . Thin controls, no zero crossing, no break
Tolerant board sensitivity	Proven on a known key: at 1–3 corrupted crib letters the exact board <i>loses</i> the true setting and a tolerance- $\geq g$ board <i>keeps</i> it, forcing 25/25 plugs correctly. Metal port cross-checks at 0/192 lane mismatches

Tolerant board specificity

Placement-dependent, not crib-length dependent. An early “ ≥ 17 edges buys tolerance 1” rule held only at offset 0 and is **withdrawn**: 18-edge menus admit 4,301 and 16,019 survivors elsewhere. Tolerance amplifies an under-determined menu rather than rescuing it, so it is a per-menu decision

Tolerance 1 against P1030680

Partial / suspended at 3/24; all three completed strongest placements are dead under right-ring coverage. Resume `--bombe-menus 0 --bombe-ring-sweep --bombe-garble-tolerance 1 --bombe-from 4`; 21 remain

Post-gap $\delta = 4$ against P1030680

Partial / suspended at 144/237; entries 1–110 and 112–144 are dead, while entry 111’s 8 raw stops yield 0 valid ≤ 10 -plug completions. Resume `--bombe-menus 0 --bombe-from 145`; middle ring remains A

Nazi Blaster 9000 (unified Verilog Mulein Future Bank)

Operational preflight and bounded production stripe DONE; 0 BREAK gates: known-key receipts agree across source RTL, post-Yosys RTL, clear Yosys-JSON, and cleartext Float TensorLUT. Runtime selected `BANK_LANES=4` at 276.035 complete receipts/s. The *Fahrenheit 261* historical 261-entry campaign executed only shell 0, identity Future 0, and settings $0 \leq s < 256$: 16/16 chunks, 6,656 receipts, zero positives, and zero BREAK gates; the other 260 Futures did not run. The separate operational inventory contains 628 Futures (314 identity + 314 post-gap $\delta = 4$). Shell 0, setting 0, and all 628 Futures completed: **628/628 chunks and 16,328 receipts**; 15 repaired/non-exact candidates yielded 0 exact hits and 0 gates. The same shell, all Futures, and settings $1 \leq s < 256$ then completed **10,048/10,048 chunks and 4,163,640 receipts**, retaining 4,184 host-verified physical candidates: 40 exact identity hits, 2,056 identity repairs, 41 exact post-gap hits, and 2,047 post-gap repairs. **0 BREAK gates; no key or decrypt**. Read-only adjudication now kills all 41 exact post-gap hits at joint ≤ 10 -plug completion, independently matching the 40 exact identity deaths: all **81 exact candidates are split-menu ghosts**. The 4,103 repairs have no declared crib-exact bar and remain deferred; broader settings and shells remain open. Bounded, not global. Cleartext, not FHE or encrypted tick rate

Soft-band Welchman stops can quarantine into Stochastic seeds (`--hybrid-quarantine`). Synthetic P1030684 block-wipe and historical P1030681 first-draft controls grade that path (escalate recovers edit ceilings, not a new key). On P1030680 itself, UEBUNG and Regenbogen soft bands produced German-looking IC without survivors under the 80% bar — coincidence, not plaintext.

6.1 A clean negative the board had not earned

The Metal bombe tabulates one upper-involution table per distinct slow-wheel (l, m) state across a menu span, and capped that at four. On overflow it wrote a zero survivor mask — the code for *eliminated* — for a lane it had never tested. Pure stepping arithmetic over every span, all 56 naval (middle, right) rotor pairs and all 676 windows (`Scripts/max_upper_audit.py`) bounds how often that fired: never at or below 25 letters, then 0.29 % of lanes at 28, 1.14 % at 30 and 5.56 % at 40. The damage is concentrated in the curated long-menu arms rather than spread thinly. Phase 16 arm 2 — forty mid-message anchors whose zero raw stops closed the Übung-push hypothesis — is 40 letters throughout, so *all forty* placements were contaminated and that negative was unearned as recorded. It has since been re-run on the corrected kernel: 40/40 dead at the board, zero raw stops over 6.387×10^{11} settings, so the cap was not concealing a key and the hypothesis stays closed on evidence. Phase 8's curated top-30, contaminated on 20 of 30 placements, likewise re-runs to 73 raw stops over 4.79×10^{11} settings — identical to the archived figures — with best board IC 0.041 / tail -4.819 and no break. In both arms the previously untested lanes yielded *zero* additional survivors. Two archival notes belong with that: the scored-completion count on the top-30 moved $0 \rightarrow 35$ purely because the head-reading gate and quarantine's soft IC floor both postdate Phase 8, the raw-stop count being the coverage-sensitive figure; and the archived top-30 log is truncated at menu 6 of 30 with no summary, so its quoted figures had never been on disk until this re-run supplied them. The remaining 300 contaminated catalog placements are retained but deprioritised: the two maximum-contamination arms independently produced zero additional survivors, while the strongest examples now also appear in the partial full-middle-ring arm. The Regenbogen queue is clean, and everything at 16–25 letters, including all of the Thetis-register work, is untouched and stands. The cap is now 8 and overflow reports *undecided*, which the host — which has no such cap — re-tests in full: a table size can no longer discard a key.

6.2 Unpinning the middle ring for 6.5–11.4×, not 26×

Before Phase 51.12, every target campaign pinned the middle Ringstellung at A, which is exact only while the middle wheel stays off its own notch in span. The gap is real (~ 8 –15 % of keys) and was priced at a flat 26×. It is cheaper, because notch tests depend on *window* position and are ring-free: if neither a lane's middle wheel nor that of its ring-A partner at window $m - \rho$ reaches a notch across the span, the middle wheel steps only on right-wheel notches — the same times in both trails — the left wheel never moves, and $\text{offset}_M = m - \rho$ is the partner's window. Same scramblers, same verdict. Ring A therefore runs in full and rings B–Z run only their notch-hitting lanes, and the union is complete coverage rather than a sample. Graded on the validated host board over 2.4×10^6 claimed-covered pairs and then through the kernel: zero verdict mismatches, and with $\sim 5\,600$ – $6\,900$ survivors per ring actually discarded, every discarded verdict is carried identically by its ring-A partner. Predicted cost is 6.5× a right-ring pass with a one-notch middle wheel and 11.4× with two. End-to-end on the blind control at full middle × right coverage — 908 544 shells, 4.152×10^{11} settings, the first run in this campaign with no residual ring gap — the break is recovered with all ten plugs in 4 307 s at 96.4 M/s: 26× the search space for 12.2× the time of the right-ring pass, because roughly 70 % of lanes exit after the trail test. Raw stops rise from 2 to 4, which is the mechanism working rather than a regression: with the middle ring unpinned the true machine has several ring-equivalent representations in the swept space. Measured basis is ~ 72 minutes per placement. Against P1030680 the arm is **suspended at 1/24**: the completed strongest placement is dead at the board over 4.152×10^{11} settings; 23 remain open. Resume `--bombe-menus 0 --bombe-middle-ring --bombe-ring-sweep --bombe-from 2`. This is the first bounded target evidence with middle ring $\neq$ A, not a verdict on the untested placements.

Two sieves were graded and rejected. The historical ten-lead board bounds the plugs a menu

has not yet forced ($26 - \text{determined} \geq 2(10 - \text{pairs})$); it is sound — the blind control still breaks — and useless, taking 847 434 raw stops to 847 414 (0.002 %) for a 4 % throughput loss, because it only bites on long strong menus that already die at the board. Separately, three quarantine flags this report and the ledger both document were never parsed, so the soft-tail arm recorded at -4.3 ran at the default -4.0 .

6.3 The crib-free path, measured

Every arm above needs a probable word, and every crib in the catalog is imported register from other nets. It is also recurrence-filtered: the miner kept only phrases carried by two or more decrypts. An internal relay audit on 2026-08-16 had already measured that 99.0 % of distinct 16–40 letter corpus windows are one-carrier phrases, yielding 2,691,455 legal placements, and emitted a 24-placement fixture. That fixture was never run, however, and the gap remained absent from this report and the public chronology. Selm Merel Wenselaers (Historian / Curator, Amsterdam / Antwerp) joined the campaign team after independently flagging the recurrence filter and supplying three specific historical leads. Her legal-placement arithmetic reproduced exactly, and her two highlighted menus are the only single-component members of their tied loop-count tier. Her work supplied the historical judgement that turned a dormant inventory into an actual experiment; HELUT’s graded machinery turned that selection into bounded measurements. These are joint campaign findings, and the combination of archival research and mechanical cryptanalysis is the point rather than a division of credit. The correction matters for interpretation: the catalog is the register *repeated across* the corpus, not its whole register. It still does not supply Thetis plaintext, which is the actual missing ingredient.

Her reported Volksliste III process—demolition parties aboard while crews transferred to Neustadt—also imposed a constraint on the machinery rather than a new crib. If two strings are linked stages of one procedure, they must be allowed to compete *independently*, and they must not be concatenated into a sentence nobody is attested to have sent. A short phrase can be historically strong and mechanically unusable at once: NEUSTADT is eight letters with no single-component placement, so the campaign’s own rules bar it from a solo sweep, and it therefore could not contribute at all. A *constellation* resolves that. It is one joint hypothesis over several independent anchors in which only the absolute (step, plaintext, ciphertext) board constraints are unioned; the anchors are never joined, and no plaintext is asserted in the gap between them. The builder is fail-closed, so a malformed, overlapping or self-enciphering anchor rejects the whole constellation rather than quietly leaving a weaker menu that still resembles the plan. Exactness became the conjunction over anchors, and the tail statistic scores only physically contiguous uncovered runs, so removing two anchors cannot manufacture an n -gram across the excision or across the gap. The independence test became constraint-disjointness, which tightens multi-menu confirmation rather than loosening it. Staged from this: FFFTTTBLEIBTBESETZT $\times$ NEUSTADT, 24 single-component rows of 27 edges and 7–9 loops, both strings repo-verbatim, topology cross-checked between an independent union-find and the board. The cost is 3.83×10^{11} settings, roughly 2.8 h.

Result: a clean negative. The arm ran to completion on 23 September 2026: 24 of 24 menus dead at the board, **zero** raw stops across 3.832×10^{11} settings at 42.5 M/s, no physically buildable board, and no soft-band candidate at all—the quarantine file was never written. Zero raw stops is a stronger shape than the campaign’s usual “many stops, none physical”: the board never produced a candidate to sieve. Under right-ring coverage with the middle ring pinned, no key places FFFTTTBLEIBTBESETZT *and* NEUSTADT at any of the tested alignments. This does not disturb the reported Volksliste III history; what it falsifies is the narrower proposition that both strings sit verbatim in this ciphertext at those offsets. The residual stays printed: a pinned middle ring leaves $\sim 12\text{--}23\%$ of keys untested across these 40–67 letter spans.

Unicity: the barrier is the estimator, not the information. This was worth computing and had not been. The full key—wheel order, Greek wheel, thin reflector, rings and position (4.151×10^{11} settings) together with the ten-plug board $\binom{26}{10 \text{ pairs}} = 26!/(6! 10! 2^{10}) \approx 1.507 \times 10^{14}$ —carries about $6.3 \times 10^{25} \approx 2^{86}$, i.e. ≈ 86 bits. German redundancy at roughly 3.2 bits per letter puts 72 letters at ≈ 230 bits, so unicity distance is near **27 letters** and the target sits at about $2.7\times$ that. Exactly one key yields German, by a margin of some 144 bits. Every negative in this report is therefore a failure of *search*, not a shortage of information: the answer provably exists inside the space being enumerated.

An independent implementation confirms that the estimator is what binds. A from-scratch Python M4, verified by reproducing every clean key in the corpus, combined with leave-one-out n -grams and *genuine* wrong-key decrypts, reproduces the standing negative: median margin -0.053 at 72 letters across orders 2–5, against the Swift pipeline’s -0.036 to -0.123 . Two unrelated implementations agreeing on a negative is stronger evidence than either alone.

One methodological trap is recorded because it nearly became a false positive. The first version of that probe reported a margin of $+1.14$ and a 100 % win rate at 72 letters, which would have overturned the central negative of this work. The defect was supplying the true plugboard to the truth *and* to every decoy: with the correct board the true setting decrypts to clean German while wrong settings are noise, so any scorer separates them perfectly. It was grading against the answer key. Requiring the board to be hill-climbed symmetrically at every candidate moved the figure from $+1.14$ to -0.044 . That single correction is the difficulty of the problem, and it is why partial exhaustion must always be applied to the truth and the decoys alike. The probe’s scope limit is equally worth stating: it trained on 4,644 letters, whereas the dense trigram table used here derives from 28.5 M. It therefore measures saturation in n -gram *order* at small sample size, not the exhaustion of language modelling as a lever.

The practical consequence is a change in ranking. The plugboard was a *daily* setting, so a second Thetis signal from 1 May would constrain the same 47-bit board with 144 letters rather than 72, and the oracle ladder already shows four correct plugs flipping the margin positive. Archival discovery therefore outranks further sweeping, including the five-to-seven day swept version of the arm above.

The measurement that matters is not the price but the geometry, and it sets the run order. A 27-edge contiguous crib spans 27 positions; these constellations span 47–67, because the anchors sit far apart. Span, not edge count, is the quantity a middle ring pinned at A fails to cover: the engine’s preflight reports $\sim 12\%$ of keys uncovered for one middle-wheel notch and $\sim 23\%$ for two across a 67-letter span, against a few percent for a compact crib. Loops are not free: the extra deduction power a joint menu buys is partly paid for in a wider blind spot, which is the residual quoted with the result above. The Ostwald–Weierud ciphertext-only attack needs none: for each candidate rotor setting it hill-climbs the plugboard and ranks the setting by the score the climb reaches. Its published reach is messages down to roughly 100 letters, with 78 the shortest broken, on three-rotor Heer traffic; P1030680 is 72 letters on four-rotor naval M4, i.e. shorter than the record on a harder machine.

The limit is not throughput, which changes how hardware should be spent. Seventy-two letters of naval German carry ~ 223 bits of redundancy, a ten-plug board is ~ 47 bits of nuisance parameter *fitted per candidate*, and the shell $\times$ position space is $\sim 10^{11}$: taking a maximum over that many overfitted candidates is a multiple-comparisons problem, and adding candidates raises the maximum of the noise. The decisive quantity is therefore the *margin* — climbed score at the true setting minus the best climbed score over wrong settings — and `--ostwald-curve` measures it at a ladder of truncation lengths over the published 1 May 1945 U-534 keys, with exhaustion applied symmetrically to true and decoy settings so the margin cannot be manufactured.

The harness refuses to report a curve unless its controls decrypt their own published plaintext under their own published key, and that preflight found something: **32 of 48 controls round-trip exactly**, the rest do not. An earlier revision read that as transcription garble corroborating

the garble hypothesis; **that reading is withdrawn**. Classified by error shape, the 17 discrepant controls have at least three causes. Most are *multi-part* messages — several exceed the ~ 320 -letter Kriegsmarine limit and were sent in parts with a fresh Grundstellung each, so end-to-end decryption under one message key correctly yields a good head and a garbage tail, which is precisely the observed shape. A few are indels or divergent transcripts. Only four controls, eight events in total, are isolated substitutions — far too thin to fit a confusion model, so any degarbler must take its prior from Girard’s direct reading of the originals (he names $U \leftrightarrow N$, $Q \leftrightarrow G$, $H \leftrightarrow F$) rather than from this corpus.

Validation is unambiguous — at 372 letters the staged climb recovers all ten plugs and 100 % of the plaintext (margin +1.64, $z = 35.6$) — which is what licenses believing the negatives. Staged scoring (IC, then bigrams, then trigrams, matching the reference implementation’s default) decisively beats the bigram-only climb previously in the repo: at 252 letters, 8/10 plugs against 0/10. The measured threshold is ~ 200 **letters**, where the median margin crosses zero; control counts above 180 are thin ($n = 3-6$), so those medians are noisy while the monotone trend is not. Ostwald’s partial exhaustion — fixing one plug and climbing the rest over all 141 candidates — removes **68 %** of the deficit at 72 letters ($-0.2389 \rightarrow -0.0754$ on identical controls) and doubles the win rate, the largest single lever found and the one the literature names.

Unseeded, the margin at 72 letters remains *negative*: the crib-free path is not yet a viable standalone arm against P1030680, our threshold is $2.8\times$ worse than the published record, and that gap is unclosed.

The lever available only here changes that picture at the target length. A Welchman stop arrives with the stecker its menu forced — 15–24 of 26 letters, i.e. 7–12 plugs — and the reference implementation documents exactly that interface (`-s`, “known plugboard connections, e.g. from running a bombe”) while requiring it to be fed by hand, because the bombe is someone else’s program. Modelled fairly, with the true setting seeded from k correct plugs and every decoy seeded from k *random* plugs, at 72 letters on ten controls the margin crosses zero at $k = 4$ (+0.130) and at $k = 8$ the attack recovers the complete key and 100 % of the plaintext on every control (+1.049, 10/10 plugs). That is below the published 78-letter record, on the harder machine — but it is *oracle-seeded*, so it assumes the seeded stop is the true key. What it establishes is narrower and still useful: a *true* bombe stop is finishable at 72 letters, so a stop the strict linguistic bar rejected for garble or short-message noise can be recovered.

That makes the climber a much better escalator than the genetic path of §Quarantine, which returned zero survivors at a $\sim 60\%$ coincidence ceiling without saying why. Escalating every soft-band stop the campaign has ever quarantined — 55 UEBUNG, 22 Regenbogen scuttle, 7 Hela, 2 from the Phase-49 re-run, 86 in total — yields no break, and in every arm the best stop scores *below* the best of twelve random settings (Δ from -0.19 to -0.80). The soft band is not a near miss; it is worse than noise under a scorer that would finish a true stop. The quarantine escalation path is therefore closed with a reason rather than a ceiling.

6.4 A diagonal board that counts contradictions

On the physical bombe a contradiction *is* current finding a second path through the diagonal board: the hypothesis short-circuits and the drum advances. Copper cannot answer “how many contradictions?”, and every faithful reimplementations inherits that, including ours, which returns **nil** on the first doubled row. Under garble this is fatal — one mis-transcribed ciphertext letter contradicts a *true* menu — so every exact-crib clean negative above is a negative about the *recorded* ciphertext rather than the transmitted one.

In silicon the question can be posed differently: how few menu edges must be deleted before a setting closes consistently? Accepting a setting when $\leq t$ deletions suffice gives a *deletion-tolerant* diagonal board, occupying the structural slot Welchman’s diagonal board did — a new board on the bombe, not a new bombe. Tolerance 0 is the historical board bit for bit.

Tolerance must *remove* edges before propagating rather than abandon them mid-flight. A first implementation did the latter and measured as an exact no-op, for two compounding reasons worth recording: by the time a conflict surfaces the bad value is already committed and has propagated through the diagonal board, and the edge visited when it surfaces is typically a *correct* edge rather than the garbled one, so blaming it leaves the real culprit in place.

Sensitivity, graded on the published P1030684 key with letters corrupted inside a 27-letter crib span: at 1, 2 and 3 garbles the exact board **loses** the true setting and a tolerance- $\geq g$ board **keeps** it, while still forcing 25/25 plug deductions correctly. Since a true stop needs only four correct plugs to flip the crib-free margin at 72 letters, the two mechanisms compose: a garbled true key survives the board *and* is finishable. The Metal port is graded against the host board rather than trusted, at **0/192** lane mismatches across 45 crib $\times$ offset $\times$ tolerance cells, with wall time per shell rising $0.03\text{ s} \rightarrow 7\text{ s}$ from tolerance 0 to 2 — confirming the extra closures are performed and not elided.

The cost is survivor inflation, and here an early reading was wrong and is withdrawn. At crib offset 0 inflation looked like a clean function of menu size: 16 edges admitted 130,787 of 456,976 lanes at tolerance 1, while ≥ 17 edges admitted exactly one, and ≥ 22 admitted one even at tolerance 2. Varying the *placement* destroys that rule: 16 edges costs nothing at offset 15, while 18-edge menus admit 4,301 and 16,019 survivors at offsets 40 and 65. Crib length does not predict inflation; the menu’s loop structure does, because tolerance spends redundancy and only a loop-rich menu has enough to spend. The placements that detonate are those already weak at tolerance 0 — tolerance *amplifies* an under-determined menu rather than rescuing it. Tolerance is therefore a per-menu decision, pre-qualified by measuring inflation against the escalator’s budget, and not a global switch.

The known-key sensitivity grade does not assert that P1030680 *is* garbled. The target tolerance-1 arm has nevertheless run to a bounded pause point: **3/24** strongest placements are complete under right-ring coverage and all three are dead at the board. It is suspended; resume `--bombe-menus 0 --bombe-ring-sweep --bombe-garble-tolerance 1 --bombe-from 4`. The remaining 21 placements are open, so this is three local negatives rather than a verdict on garble generally.

Indels are the board’s *second* mechanism rather than an omission: Girard records an entire four-letter group (HMHY) present on one transcript and absent from the other. That frame shift is handled by re-indexing the menu and shifting step numbers — a geometry change on the exact board, with no survivor inflation or pre-qualification gate. Against P1030680, the post-gap $\delta = 4$ arm is **suspended at 144/237**: entries 1–110 and 112–144 are dead; entry 111’s 8 raw stops produce no valid ≤ 10 -plug completion. Resume `--bombe-menus 0 --bombe-from 145`; 93 hypotheses remain and the middle ring is still A. No physical candidate or break has resulted.

A production Verilog Future Bank unifies explicit Future geometry, exact-first one-edge repair, a shared 80×26 trail, and held receipts. On P1030684 controls, complete receipts agree across source RTL, post-Yosys RTL, clear Yosys-JSON simulation, and cleartext Float TensorLUT, including held-valid backpressure.

Identical complete-receipt benchmarks then compared bank widths 1/2/4/8/16 on 16 settings $\times$ 16 jobs $\times$ five repetitions. All widths produced digest `f1a64-17435188996435da`; median rates were 233.777350172, 264.718510307, **276.035471632**, 273.889660462, and 230.780129610 receipts/s. Thus `BANK_LANES=4` was selected by runtime rather than synthesis size. Its graph has 189,032 LUT6, 2,772 DFF, 205,177 wires, and 28 levels; the selection receipt is `logs/mulein-future-tensorlut-selection.json`.

The protocol-v3 target runner binds one immutable manifest byte snapshot, the netlist, graph, width, protocol, and exact shell/Future/setting plan into its run identity. Resume admits only an exact ordered plan prefix. Every synchronized chunk stores the canonical per-job receipt projection, permitting digest recomputation and derivation of hit and BREAK-gate state rather than trust in summary fields. Persisted hits are host/scoring-replayed; a prior valid gate halts

before evaluator submission. Semantic row drift, sparse or reordered chunks, candidate deletion, and records after a gate fail closed. Atomic non-truncating open and a retained exclusive lock enforce one writer through final synchronization. Metal status, tags, seeds, held receipts, consume/rearm, and timeout are also checked. Only exact identity/no-drop positives may enter the existing discriminator; repaired and post-gap positives remain non-BREAK-eligible until their correction or full-message geometry is independently replayed.

The historical *Fahrenheit 261* campaign stages 261 menu/Future entries, but its bounded launch covered only shell 0 B/beta/IV-III-VIII/AAAA, identity Future 0, and settings $0 \leq s < 256$. Protocol-v3 run sha256-e6dc10d4...2fc1360 completed **16/16 synchronized chunks and 6,656/6,656 canonical receipt projections, with zero hardware positives and zero BREAK gates**; identical resume revalidated all 16 rows and appended nothing. Durable receipt: logs/p1030680-mulein-unified-smoke-v3.jsonl, SHA-256 687d0838...2b1f7e7d. This remains one clean historical local negative, not coverage of that manifest's other 260 Futures.

Nazi Blaster 9000, the formal name of the unified P1030680 Mulein search machine and overall campaign, then introduced a separate operational/scuttle inventory rather than rewriting the Fahrenheit 261 artifact: 628 Futures (314 identity + 314 post-gap $\delta = 4$), fingerprint fnv1a64-616326e94036a97d. Shell 0, setting 0, and all 628 Futures completed **628/628 chunks and 16,328 checked receipts**. Fifteen host-replayed physical candidates survived—5 identity-family one-edge repairs and 10 post-gap one-edge repairs—but every candidate was non-exact. Therefore the grade is **0 exact hits and 0 BREAK gates**, not “zero positives.” Repaired candidates remain non-BREAK-eligible without explicit correction replay; post-gap candidates require geometry-aware full-message replay. They do not establish that the ciphertext is garbled. The durable receipt is logs/p1030680-mulein-operational-preflight-v3.jsonl, SHA-256 55a68266...9a17f993.

The bounded production stripe is **DONE as bounded coverage**: the same shell, all 628 operational Futures, and settings $1 \leq s < 256$. From 2026-08-20T20:44:18Z through 2026-08-21T01:54:45Z, protocol-v3 run sha256-efb770bc...592cc0 wrote **10,048/10,048 new synchronized chunks, 0 skipped, and 4,163,640/4,163,640 canonical receipt projections**. The durable receipt is the settings-000001-000256-v3 production ledger under logs/; its SHA-256 is c516238d...b1e7e52f2.

Host replay retained **4,184 physical candidates**, all verified: 40 exact/no-drop identity hits killed by ≤ 10 -plug completion, 2,056 one-edge identity repairs, 41 exact/no-drop post-gap $\delta = 4$ hits, and 2,047 one-edge post-gap repairs. There are **0 row BREAK gates, 0 candidate BREAK gates, no scored plaintext, no key, and no decrypt**. This is not a zero-positive result or a clean global negative. A translation audit paired 1,724 identity/post-gap candidates at a four-setting-lane shift; over 82% of each family therefore resamples the same closure/live state in this stripe. That result does not eliminate a missing group at another physical boundary, but makes blind widening and same-shape leading $\delta = 6/8$ stripes low-information.

The next justified action is geometry-aware completion and sparse full-message scoring of the 41 exact post-gap hits, followed by deduplicated explicit corrected-letter replay of the highest-constraint identity repairs. Repaired post-gap hits rank third; the 40 exact identity hits need no rerun because completion already killed them. Operational/scuttle hypotheses ran first as a working search prior, not historical proof. Settings $256 \leq s < 456,976$, other shells, and the catalog, middle-ring, tolerance, and legacy-indel remainders remain open. The rates, receipts, and fitness are cleartext Float TensorLUT/Metal—not encrypted tick rate, not FHE, not a key, and not a decrypt.

7 Parallel Track: TensorLUT

Independently of the Boolean catalog, HELUT ships a continuous-to-discrete compiler: Yosys LUT6 cells become 64-wide floating INIT tensors, Metal evaluates a multilinear stream, λ

squeezes toward binary, and an emitter writes gate-level Verilog. Unmutated `enigma.m4` locks $F_{\text{crypto}} = 0$ (925 LUT6 + 49 DFFs baseline). Full INIT cold-start and stecker-cone melts *shatter* under λ ; the live arm freezes known-good silicon and evolves a ≤ 10 -pair stecker involution by construction (blind three-pair rediscovery PASS; claims C8–C9). Six formal lemmas of the continuous–discrete loop hold in-repo (claim C19; `directives/tensorlut-theorem.md`). Melt–freeze–snap on a separable Boolean interpolant is claim C44; a 2-LUT cascade melt that emits Verilog and matches 8 Boolean corners is claim C48 (dual interpolants allowed). This grades generative hardware synthesis and reciprocal genotype search. It does **not** invent P1030680’s missing plaintext (N7, H6).

8 Scope Relative to the HELUT Stack

HELUT also graduates a torus FHE path (LWE/GLWE, GGSW blind-rotate, certificates) documented in `paper/helut.tex` and `directives/claim-sheet.md`. That path is out of scope for campaign fitness: P1030680 search remains cleartext Metal batch. Encrypted SING envelopes and hardness calibration are stack claims (C4–C7, H1–H3), not evidence that this message is broken.

9 Conclusion

P1030680 remains unbroken. The contribution of this campaign is a graded, reproducible Boolean and stochastic laboratory on Apple Silicon; archival elimination of wrong nets and many wrong wedges; scoring and ghost-control fixes that make short-M4 search honest; the benchmark-selected *Nazi Blaster 9000* unified Verilog-to-TensorLUT machine; and durable bounded receipts for the *Fahrenheit 261* historical slice, the all-Future operational preflight, and the completed settings-1...255 production stripe. That stripe retained 4,184 host-verified physical candidates but fired zero BREAK gates and supplied no scored plaintext, key, or decrypt; repaired and post-gap candidates remain outside the gate pending explicit replay. A parallel TensorLUT arm learns reciprocal structure without claiming a decrypt. Geometry-aware candidate replay is now complete for every *exact* candidate, and it corrected this campaign’s own diagnosis. Counting the durable ledger rather than quoting it gave **41** exact/no-drop post-gap candidates, **40** exact identity candidates already killed by plug completion, and **4,103** repaired. A read-only adjudicator then graded the 41 with no GPU, because the ledger already persists shell, setting, seed, Future identity, and repair provenance; it writes its own ledger, since the campaign ledger re-derives every candidate and compares by equality, so mutation would invalidate all 4,184 rows. Host replay verified **41/41**, and **all 41 return no-consistent-board**: no ≤ 10 -plug involution satisfies the whole menu, so they never reach the language bar and the missing geometry-aware decrypt was never the binding constraint. Their six source Futures carry **two to five components** on 16–21 edges, one a zero-loop tree, making them split-menu ghosts of the shape documented since Phase 4 — and the campaign had already killed all 40 identity siblings by the same sieve through an independent path. **All 81 exact candidates in that stripe are ghosts.** A P1030684 calibration control gates the scorer and additionally proves the sparse order-3 mean reproduces the dense trigram score *bit-exactly* on hole-free geometry, so the $-3.600 / 0.055$ bars transfer in scale after all; what does not transfer is the dense crib-excising *sample*. The **4,103 repaired candidates are deferred**: a repaired hit is crib-exact-minus- k by definition, so they need a declared bar and human sign-off rather than more code. The actionable lesson is about inventory, not compute: the next manifest should gate on connectivity rather than crib length, which would have excluded all six of those Futures before a single setting ran. The completed stripe is bounded evidence, not a global negative; broader settings, shells, the middle ring, and the printed legacy remainders remain open.

Acknowledgements

Frode Weierud and Olaf Ostwald's published ciphertext-only work is the basis of the crib-free arm measured above. The U-534 corpus, the transcriptions, and the sister-message degarbling that motivates the tolerant board are the work of Michael Hörenberg and Jean-François Girard; this campaign is downstream of their archival effort and would not exist without it.

Selm Merel Wenselaers (Historian / Curator, Amsterdam / Antwerp), a deeply valued member of the campaign team, joined the work after independently identifying the recurrence filter and supplying three verbatim-attested, historically motivated one-carrier leads. Her 15-of-44 legal-placement calculation reproduced exactly; her two highlighted placements were also the only single-component menus among the four tied at the highest loop count. Her archival judgement moved a dormant internal surface into an actual campaign and public view. The resulting work is joint: Selm's historical and archival research defines evidence-backed questions, and HELUT's graded machinery turns them into controlled measurements. The campaign is stronger for combining both disciplines, and the findings belong to that collaboration.

For chronology, HELUT's `Scripts/relay_crib_mine.py` had enumerated the general 99.0 % hapax surface on 2026-08-16, but its 24-placement fixture had never been run or reported. Selm's independent diagnosis, exact leads, placement arithmetic, and menu-selection insight supplied the historical prioritisation that made the surface actionable. She now leads the campaign's historical and archival arm, including primary-source and collection enquiries, while remaining part of the decision loop before new computational tranches begin.

Naming collaborators records provenance; it does not imply that any person named above has claimed a decrypt of P1030680.